

Lab 3: Data Preparation

CPE232 Data Models

Owner : 66070501043 - Phoorin Chinphuad

[1] Reviews on Pandas

1.1) Discover

- methods to explore and understand your DataFrame

```
In [159... import pandas as pd

df = pd.read_csv('./sources/lab/nss15.csv')
```

```
In [160... # see the shape of the dataframe
print(df.shape)
```

(334839, 12)

```
In [161... # seeing the summary of the dataframe
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 334839 entries, 0 to 334838
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   caseNumber            334839 non-null int64  
 1   treatmentDate         334839 non-null object 
 2   statWeight            334839 non-null float64 
 3   stratum               334839 non-null object 
 4   age                   334839 non-null int64  
 5   sex                   334837 non-null object 
 6   race                  205014 non-null object 
 7   diagnosis             334839 non-null int64  
 8   bodyPart              334839 non-null int64  
 9   disposition           334839 non-null int64  
10   location              334839 non-null int64  
11   product               334839 non-null int64  
dtypes: float64(1), int64(7), object(4)
memory usage: 30.7+ MB
None
```

```
In [162... # seeing the stats of the column in dataframe
print(df.describe())
```

	caseNumber	statWeight	age	diagnosis \
count	3.348390e+05	334839.000000	334839.000000	334839.000000
mean	1.510271e+08	39.343028	31.385451	60.154591
std	1.720330e+06	34.142933	26.105098	6.170699
min	1.501032e+08	4.965500	0.000000	41.000000
25%	1.504405e+08	15.059100	10.000000	57.000000
50%	1.507358e+08	15.776200	23.000000	59.000000
75%	1.510231e+08	74.881300	51.000000	64.000000
max	1.603418e+08	97.923900	107.000000	74.000000

	bodyPart	disposition	location	product
count	334839.000000	334839.000000	334839.000000	334839.000000
mean	64.374192	1.307930	2.485451	2098.900854
std	24.002331	0.977627	3.217617	1332.222670
min	0.000000	1.000000	0.000000	106.000000
25%	35.000000	1.000000	0.000000	1211.000000
50%	75.000000	1.000000	1.000000	1807.000000
75%	82.000000	1.000000	5.000000	3265.000000
max	94.000000	9.000000	9.000000	5555.000000

In [163... `# seeing the first 5 rows of the dataframe`
`print(df.head())`

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race \
0	150733174	7/11/2015	15.7762	V	5	Male	NaN
1	150734723	7/6/2015	83.2157	S	36	Male	White
2	150817487	8/2/2015	74.8813	L	20	Female	NaN
3	150717776	6/26/2015	15.7762	V	61	Male	NaN
4	150721694	7/4/2015	74.8813	L	88	Female	Other

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893

In [164... `# seeing the last 5 rows of the dataframe`
`print(df.tail())`

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race \
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN
334835	150733393	7/11/2015	5.6748	C	3	Female	Black
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN
334837	150823002	8/8/2015	97.9239	M	38	Female	White
334838	150723074	6/20/2015	49.2646	M	5	Female	White

	diagnosis	bodyPart	disposition	location	product
334834	59	76	1	1	1864
334835	68	85	1	0	1931
334836	71	79	1	0	3250
334837	59	82	1	1	464
334838	57	34	1	9	3273

In [165... `# seeing the list of columns in the dataframe`
`print(df.columns)`

```
Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'age', 'sex',
      'race', 'diagnosis', 'bodyPart', 'disposition', 'location', 'product'],
      dtype='object')
```

1.2) Selecting variables

- select specific columns from the DataFrame to create a new DataFrame with only those columns

```
In [166... df['age']
```

```
Out[166... 0      5
1     36
2     20
3     61
4     88
..
334834    7
334835    3
334836   38
334837   38
334838    5
Name: age, Length: 334839, dtype: int64
```

```
In [167... df['age'].head()
```

```
Out[167... 0      5
1     36
2     20
3     61
4     88
Name: age, dtype: int64
```

```
In [168... df[['caseNumber', 'age']]
```

```
Out[168...
```

	caseNumber	age
0	150733174	5
1	150734723	36
2	150817487	20
3	150717776	61
4	150721694	88
...
334834	150739278	7
334835	150733393	3
334836	150819286	38
334837	150823002	38
334838	150723074	5

334839 rows × 2 columns

```
In [169... # select columns based on the data type
df.select_dtypes(include=['number'])
```

Out[169...

	caseNumber	statWeight	age	diagnosis	bodyPart	disposition	location	product
0	150733174	15.7762	5	57	33	1	9	1267
1	150734723	83.2157	36	57	34	1	1	1439
2	150817487	74.8813	20	71	94	1	0	3274
3	150717776	15.7762	61	71	35	1	0	611
4	150721694	74.8813	88	62	75	1	0	1893
...
334834	150739278	15.0591	7	59	76	1	1	1864
334835	150733393	5.6748	3	68	85	1	0	1931
334836	150819286	15.7762	38	71	79	1	0	3250
334837	150823002	97.9239	38	59	82	1	1	464
334838	150723074	49.2646	5	57	34	1	9	3273

334839 rows × 8 columns

In [170...

```
# select row by .loc
df.loc[0]
```

Out[170...

caseNumber 150733174
treatmentDate 7/11/2015
statWeight 15.7762
stratum V
age 5
sex Male
race NaN
diagnosis 57
bodyPart 33
disposition 1
location 9
product 1267
Name: 0, dtype: object

In [171...

```
# select column by .loc
df.loc[:,6,'treatmentDate':'diagnosis']
```

Out[171...

	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
0	7/11/2015	15.7762	V	5	Male	NaN	57
1	7/6/2015	83.2157	S	36	Male	White	57
2	8/2/2015	74.8813	L	20	Female	NaN	71
3	6/26/2015	15.7762	V	61	Male	NaN	71
4	7/4/2015	74.8813	L	88	Female	Other	62
5	7/2/2015	5.6748	C	1	Female	White	71
6	6/8/2015	15.7762	V	25	Male	Black	51

In [172...

```
df.loc[df['age']>80, ['treatmentDate', 'age']]
```

Out[172...

	treatmentDate	age
4	7/4/2015	88
8	7/16/2015	98
39	5/3/2015	88
46	4/15/2015	91
63	1/12/2015	97
...
334701	4/27/2015	86
334784	7/7/2015	82
334785	7/11/2015	86
334815	10/28/2015	85
334819	1/13/2015	85

20422 rows × 2 columns

In [173...

```
# select row by .iloc
df.iloc[0:5]
```

Out[173...

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	dis
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	33	
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	34	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	94	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	35	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	75	

In [174...

```
# select column by .iloc
df.iloc[:,[0,1,2,3,4]]
```

Out[174...

	caseNumber	treatmentDate	statWeight	stratum	age
0	150733174	7/11/2015	15.7762	V	5
1	150734723	7/6/2015	83.2157	S	36
2	150817487	8/2/2015	74.8813	L	20
3	150717776	6/26/2015	15.7762	V	61
4	150721694	7/4/2015	74.8813	L	88
...
334834	150739278	5/31/2015	15.0591	V	7
334835	150733393	7/11/2015	5.6748	C	3
334836	150819286	7/24/2015	15.7762	V	38
334837	150823002	8/8/2015	97.9239	M	38
334838	150723074	6/20/2015	49.2646	M	5

334839 rows × 5 columns

1.3) Filtering the data

In [175...

```
# filter rows based on the condition
df[df['age'] > 50]
```

Out[175...

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPar
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	3
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	7
7	150704114	6/14/2015	83.2157	S	53	Male	White	57	3
8	150736558	7/16/2015	83.2157	S	98	Male	Black	59	7
16	150901411	8/27/2015	83.2157	S	65	Female	White	59	8
...
334811	150702215	6/27/2015	15.7762	V	51	Female	NaN	53	8
334815	151100368	10/28/2015	83.2157	S	85	Female	NaN	57	8
334819	150528367	1/13/2015	49.2646	M	85	Female	NaN	57	7
334826	150648619	6/17/2015	15.7762	V	52	Female	White	64	3
334829	150633526	4/4/2015	49.2646	M	51	Female	NaN	56	9

85235 rows × 12 columns



In [176...

```
# filter coloum based on column name
df.filter(like='age')
```

Out[176...

age	
0	5
1	36
2	20
3	61
4	88
...	...
334834	7
334835	3
334836	38
334837	38
334838	5

334839 rows × 1 columns

1.4) Sorting

- Sort the DataFrame by its index based on column

In [177...

```
# sort the dataframe based on column name and ascending order
df.sort_values(by='statWeight', ascending=False)
```

Out[177...

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPar
59985	151114128	11/1/2015	97.9239	M	17	Male	Black	64	3
239709	150809041	8/2/2015	97.9239	M	34	Female	White	57	8
239711	151018719	10/5/2015	97.9239	M	72	Female	Black	59	3
239659	151138100	10/25/2015	97.9239	M	24	Male	NaN	64	3
239696	150903084	8/25/2015	97.9239	M	75	Male	NaN	53	7
...
123586	151201944	11/28/2015	4.9655	C	11	Female	White	56	3
123589	151226169	12/8/2015	4.9655	C	5	Male	Black	64	3
138289	151132404	11/8/2015	4.9655	C	6	Female	White	57	7
138295	151136771	11/14/2015	4.9655	C	5	Female	Black	59	7
138235	151150939	11/19/2015	4.9655	C	14	Female	White	64	9

334839 rows × 12 columns



In [178...

```
# sort the index of the dataframe
df.sort_index()
```

Out[178...

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPar
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	3
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	3
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	9
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	3
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	7
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59	7
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68	8
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71	7
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59	8
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57	3

334839 rows × 12 columns



1.5) Add/Remove

- This section shows how to manipulate the DataFrame's structure

In [179...

```
# Dropping the column
df.drop(columns=['disposition'])
```

Out[179...

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPar
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	3
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	3
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	9
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	3
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	7
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59	7
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68	8
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71	7
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59	8
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57	3

334839 rows × 11 columns




In [180...

```
# Adding column and create into a new column
df.assign(new_column=df['diagnosis'] + df['bodyPart'])
```


Out[180...

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPar
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	3
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	3.
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	9.
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	3
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	7
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59	7
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68	8
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71	7
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59	8
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57	3.

334839 rows × 13 columns



In [181...

```
# Removing the column and assigning it to a new variable
ages = df.pop('age')
```

1.6) Clean missing

- to remove rows with missing values or replace missing values with a specified value


In [182...

```
# replacing the missing values with a specified value
df.fillna(value=0)
```

Out[182...

	caseNumber	treatmentDate	statWeight	stratum	sex	race	diagnosis	bodyPart	dis
0	150733174	7/11/2015	15.7762	V	Male	0	57	33	
1	150734723	7/6/2015	83.2157	S	Male	White	57	34	
2	150817487	8/2/2015	74.8813	L	Female	0	71	94	
3	150717776	6/26/2015	15.7762	V	Male	0	71	35	
4	150721694	7/4/2015	74.8813	L	Female	Other	62	75	
...	
334834	150739278	5/31/2015	15.0591	V	Male	0	59	76	
334835	150733393	7/11/2015	5.6748	C	Female	Black	68	85	
334836	150819286	7/24/2015	15.7762	V	Male	0	71	79	
334837	150823002	8/8/2015	97.9239	M	Female	White	59	82	
334838	150723074	6/20/2015	49.2646	M	Female	White	57	34	

334839 rows × 11 columns



In [183...

```
# Remove the rows with missing values
df.dropna()
```

Out[183...

	caseNumber	treatmentDate	statWeight	stratum	sex	race	diagnosis	bodyPart	dis
1	150734723	7/6/2015	83.2157	S	Male	White	57	34	
4	150721694	7/4/2015	74.8813	L	Female	Other	62	75	
5	150721815	7/2/2015	5.6748	C	Female	White	71	76	
6	150713483	6/8/2015	15.7762	V	Male	Black	51	33	
7	150704114	6/14/2015	83.2157	S	Male	White	57	30	
...
334830	150628863	6/8/2015	15.7762	V	Female	White	64	79	
334831	150607637	5/22/2015	5.6748	C	Female	Black	59	94	
334835	150733393	7/11/2015	5.6748	C	Female	Black	68	85	
334837	150823002	8/8/2015	97.9239	M	Female	White	59	82	
334838	150723074	6/20/2015	49.2646	M	Female	White	57	34	

205014 rows × 11 columns



[2] Data Cleaning and Preparation

.isnull, .dropna, .fillna

2.1) checking

In [184...

```
df.columns
```

Out[184...

```
Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'sex', 'race',  
      'diagnosis', 'bodyPart', 'disposition', 'location', 'product'],  
      dtype='object')
```

In [185...

```
# isnull checking
df.isnull().sum()
```

Out[185...

```
caseNumber      0  
treatmentDate   0  
statWeight      0  
stratum         0  
sex             2  
race          129825  
diagnosis       0  
bodyPart        0  
disposition     0  
location        0  
product         0  
dtype: int64
```

```
In [186... # percentage of missing values for the race
df.race.isnull().sum()/df.shape[0]*100
```

```
Out[186... np.float64(38.772365226272925)
```

```
In [187... df.shape[0]
```


```
Out[187... 334839
```

2.2) Drop column

```
In [188... # remove column by using
df = df.drop(columns=['race'])
```

```
In [189... df.head()
```

```
Out[189...      caseNumber  treatmentDate  statWeight  stratum    sex  diagnosis  bodyPart  disposition  loca
0    150733174      7/11/2015     15.7762      V    Male      57      33      1
1    150734723      7/6/2015     83.2157      S    Male      57      34      1
2    150817487      8/2/2015     74.8813      L  Female      71      94      1
3    150717776      6/26/2015     15.7762      V    Male      71      35      1
4    150721694      7/4/2015     74.8813      L  Female      62      75      1
```



2.3) Data imputation

```
In [190... # fillna
df['age'] = df['age'].fillna(df['age'].median())
```

```

-----
KeyError                                Traceback (most recent call last)
File c:\Users\Feen_Phoorin\AppData\Local\Programs\Python\Python313\Lib\site-packages\pandas\
\core\indexes\base.py:3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas\_libs\hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.PyObjectHashT
able.get_item()

File pandas\_libs\hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.PyObjectHashT
able.get_item()

KeyError: 'age'

```

The above exception was the direct cause of the following exception:

```

KeyError                                Traceback (most recent call last)
Cell In[190], line 2
      1 # fillna
----> 2 df['age'] = df['age'].fillna(df['age'].median())

File c:\Users\Feen_Phoorin\AppData\Local\Programs\Python\Python313\Lib\site-packages\pandas\
\core\frame.py:4102, in DataFrame._getitem__(self, key)
    4100 if self.columns.nlevels > 1:
    4101     return self._getitem_multilevel(key)
-> 4102 indexer = self.columns.get_loc(key)
    4103 if is_integer(indexer):
    4104     indexer = [indexer]

File c:\Users\Feen_Phoorin\AppData\Local\Programs\Python\Python313\Lib\site-packages\pandas\
\core\indexes\base.py:3812, in Index.get_loc(self, key)
    3807 if isinstance(casted_key, slice) or (
    3808     isinstance(casted_key, abc.Iterable)
    3809     and any(isinstance(x, slice) for x in casted_key)
    3810 ):
    3811     raise InvalidIndexError(key)
-> 3812 raise KeyError(key) from err
    3813 except TypeError:
    3814     # If we have a listlike key, _check_indexing_error will raise
    3815     # InvalidIndexError. Otherwise we fall through and re-raise
    3816     # the TypeError.
    3817     self._check_indexing_error(key)

KeyError: 'age'

```

[Q1] From the above cell, Why it showing an error?

Ans: **Because the line `ages = df.pop('age')` removes the 'age' column from df, so when I try to fill missing values in `df['age']`, the column no longer exists.**

[Q2] Fix the error from Q1 problem.

```

In [191... # [Q2]

# hint: see the cell that run `df.pop()`
df["age"] = ages

```

```
# fillna again
df['age'] = df['age'].fillna(df['age'].median())

df.head()
```

Out[191...

	caseNumber	treatmentDate	statWeight	stratum	sex	diagnosis	bodyPart	disposition	loc
0	150733174	7/11/2015	15.7762	V	Male	57	33	1	
1	150734723	7/6/2015	83.2157	S	Male	57	34	1	
2	150817487	8/2/2015	74.8813	L	Female	71	94	1	
3	150717776	6/26/2015	15.7762	V	Male	71	35	1	
4	150721694	7/4/2015	74.8813	L	Female	62	75	1	

2.4) Drop row that have missing value

In [192...

```
# remove column by using .dropna()
df = df.dropna()
```

In [193...

```
df.isnull().sum()
```

Out[193...

```
caseNumber      0
treatmentDate   0
statWeight      0
stratum         0
sex             0
diagnosis       0
bodyPart        0
disposition     0
location        0
product         0
age             0
dtype: int64
```

Datetime

2.5) Working with the datetime format

In [194...

```
df["treatmentDate"] = pd.to_datetime(df["treatmentDate"], format="%m/%d/%Y")
```

C:\Users\Feen_Phoorin\AppData\Local\Temp\ipykernel_19824\3208943844.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df["treatmentDate"] = pd.to_datetime(df["treatmentDate"], format="%m/%d/%Y")
```

In [195...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 334837 entries, 0 to 334838
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   caseNumber      334837 non-null  int64
1   treatmentDate   334837 non-null  datetime64[ns]
2   statWeight      334837 non-null  float64
3   stratum         334837 non-null  object
4   sex             334837 non-null  object
5   diagnosis       334837 non-null  int64
6   bodyPart        334837 non-null  int64
7   disposition     334837 non-null  int64
8   location        334837 non-null  int64
9   product         334837 non-null  int64
10  age             334837 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(7), object(2)
memory usage: 30.7+ MB
```

```
In [196... df['Year'] = df['treatmentDate'].dt.year
```

C:\Users\Feen_Phoorin\AppData\Local\Temp\ipykernel_19824\1686165144.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Year'] = df['treatmentDate'].dt.year
```

```
In [197... df['Month'] = df['treatmentDate'].dt.month
```

C:\Users\Feen_Phoorin\AppData\Local\Temp\ipykernel_19824\404848564.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Month'] = df['treatmentDate'].dt.month
```

```
In [198... df.head()
```

```
Out[198...
   caseNumber  treatmentDate  statWeight  stratum  sex  diagnosis  bodyPart  disposition  loc:
0    150733174    2015-07-11    15.7762      V    Male      57      33      1
1    150734723    2015-07-06    83.2157      S    Male      57      34      1
2    150817487    2015-08-02    74.8813      L  Female      71      94      1
3    150717776    2015-06-26    15.7762      V    Male      71      35      1
4    150721694    2015-07-04    74.8813      L  Female      62      75      1
```

[Q3] Can you change the format to DD/MM/YYYY? Show your work.

```
In [199... # change the format to DD/MM/YYYY
df["treatmentDate"] = df["treatmentDate"].dt.strftime("%d/%m/%Y")
df
```

```
C:\Users\Feen_Phoorin\AppData\Local\Temp\ipykernel_19824\2894307848.py:2: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
```

```
df["treatmentDate"] = df["treatmentDate"].dt.strftime("%d/%m/%Y")
```

Out[199...

	caseNumber	treatmentDate	statWeight	stratum	sex	diagnosis	bodyPart	disposition
0	150733174	11/07/2015	15.7762	V	Male	57	33	
1	150734723	06/07/2015	83.2157	S	Male	57	34	
2	150817487	02/08/2015	74.8813	L	Female	71	94	
3	150717776	26/06/2015	15.7762	V	Male	71	35	
4	150721694	04/07/2015	74.8813	L	Female	62	75	
...
334834	150739278	31/05/2015	15.0591	V	Male	59	76	
334835	150733393	11/07/2015	5.6748	C	Female	68	85	
334836	150819286	24/07/2015	15.7762	V	Male	71	79	
334837	150823002	08/08/2015	97.9239	M	Female	59	82	
334838	150723074	20/06/2015	49.2646	M	Female	57	34	

334837 rows × 13 columns



Combine Dataframe by .merge and .concat

2.6 Merge

In [280...

```
import pandas as pd
```

```
superstore_order = pd.read_csv('./sources/lab/Superstore/superstore_order.csv')
superstore_people = pd.read_csv('./sources/lab/Superstore/superstore_people.csv')
superstore_return = pd.read_csv('./sources/lab/Superstore/superstore_return.csv')
```

In [281...

```
superstore_order.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8880 entries, 0 to 8879
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row ID                8880 non-null  int64
1   Order ID              8880 non-null  object
2   Order Date            8880 non-null  object
3   Ship Date             8880 non-null  object
4   Ship Mode             8880 non-null  object
5   Customer ID           8880 non-null  object
6   Customer Name         8880 non-null  object
7   Segment              8880 non-null  object
8   Country               8880 non-null  object
9   City                 8880 non-null  object
10  State                8880 non-null  object
11  Postal Code           8880 non-null  int64
12  Region               8880 non-null  object
13  Product ID            8880 non-null  object
14  Category              8880 non-null  object
15  Sub-Category          8880 non-null  object
16  Product Name          8880 non-null  object
17  Sales                 8880 non-null  float64
18  Quantity              8880 non-null  int64
19  Discount              8880 non-null  float64
20  Profit               8880 non-null  float64
dtypes: float64(3), int64(3), object(15)
memory usage: 1.4+ MB
```

In [282... `superstore_people.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Person  4 non-null      object
1   Region  4 non-null      object
dtypes: object(2)
memory usage: 196.0+ bytes
```

In [283... `superstore_return.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 296 entries, 0 to 295
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Returned    296 non-null   object
1   Order ID    296 non-null   object
dtypes: object(2)
memory usage: 4.8+ KB
```

In [284... `superstore_order.merge(superstore_return[superstore_return["Returned"]=="Yes"],`
`on="Order ID" ,`
`how="inner")\`
`[["Customer ID", "Returned"]]\`
`.drop_duplicates()`

Out[284...

	Customer ID	Returned
0	ZD-21925	Yes
3	TB-21055	Yes
10	JS-15685	Yes
13	LC-16885	Yes
20	BS-11755	Yes
...
688	ED-13885	Yes
689	TS-21205	Yes
696	MF-17665	Yes
702	SH-19975	Yes
705	RB-19435	Yes

222 rows × 2 columns

[Q4] What does the argument `how="inner"` do?

Ans: **Keeps only the rows with matching values in both DataFrames based on the specified on = column**

[Q5] In your opinion, what information that the result above conveys?

Ans: **The result above conveys customer return behavior by listing customers who have returned at least one product.**

More merging...


In [285...

```
superstore_order.merge(superstore_return,  
on="Order ID" ,  
how="inner")
```

Out[285...

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	
0	19	CA-2014-143336	27/08/2014	01/09/2014	Second Class	ZD-21925	Zuschuss Donatelli	Consumer	United States	Fr
1	20	CA-2014-143336	27/08/2014	01/09/2014	Second Class	ZD-21925	Zuschuss Donatelli	Consumer	United States	Fr
2	21	CA-2014-143336	27/08/2014	01/09/2014	Second Class	ZD-21925	Zuschuss Donatelli	Consumer	United States	Fr
3	56	CA-2016-111682	17/06/2016	18/06/2016	First Class	TB-21055	Ted Butterfield	Consumer	United States	
4	57	CA-2016-111682	17/06/2016	18/06/2016	First Class	TB-21055	Ted Butterfield	Consumer	United States	
...	
702	8870	CA-2017-101805	01/12/2017	06/12/2017	Standard Class	SH-19975	Sally Hughsby	Corporate	United States	
703	8871	CA-2017-101805	01/12/2017	06/12/2017	Standard Class	SH-19975	Sally Hughsby	Corporate	United States	
704	8872	CA-2017-101805	01/12/2017	06/12/2017	Standard Class	SH-19975	Sally Hughsby	Corporate	United States	
705	8873	US-2014-105137	10/10/2014	10/10/2014	Same Day	RB-19435	Richard Bierner	Consumer	United States	Co
706	8874	US-2014-105137	10/10/2014	10/10/2014	Same Day	RB-19435	Richard Bierner	Consumer	United States	Co

707 rows × 22 columns



2.7) Concatenate


In [286...

```
pd.concat([superstore_order, superstore_people], axis=1, join='inner')
```

Out[286...

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Hend
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Hend
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	An
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean ODonnell	Consumer	United States	Laude

4 rows × 23 columns



[Q6] What is the difference between `inner` and `outer` on parameter `join` in `pd.concat` ?

Ans:

- **Inner Join:** Included only the rows with matching indices in both DataFrames.
- **Outer Join:** Includes all rows in both DataFrames, filling missing values with NaN.

Groupby

In [287...

```
superstore_order.groupby(['Segment', 'Ship Mode'])[['Sales', 'Quantity', 'Discount', 'Profit']]
```

Out[287...

		Sales	Quantity	Discount	Profit
Segment	Ship Mode				
Consumer	First Class	138594.9328	2455	110.29	18953.7264
	Same Day	53660.6340	1001	43.85	8555.7193
	Second Class	203605.6822	3489	127.29	24701.9148
	Standard Class	627061.3262	10430	443.05	68864.9892
Corporate	First Class	97720.1209	1670	73.07	12660.2526
	Same Day	41716.5550	366	14.50	1120.9222
	Second Class	130759.9288	2027	71.47	15582.1762
	Standard Class	359359.2109	6203	262.82	49832.6780
Home Office	First Class	76743.8674	924	39.82	11829.8821
	Same Day	20968.5170	343	12.50	3909.3442
	Second Class	77175.1080	1148	37.80	12785.8953
	Standard Class	218325.9795	3595	142.14	27298.5786

[Q7] Describe an information that the result above conveys?

Ans: **The result above conveys a summation of the sales activity, profit, quantity sold, and discounting patterns based on the customer segments and shipping modes.**

In [288...

```
superstore_order["Profit Ratio"] = superstore_order["Profit"]/superstore_order["Sales"]
superstore_order[['Profit', 'Sales', 'Profit Ratio']]
```

Out[288...

	Profit	Sales	Profit Ratio
0	41.9136	261.9600	0.1600
1	219.5820	731.9400	0.3000
2	6.8714	14.6200	0.4700
3	-383.0310	957.5775	-0.4000
4	2.5164	22.3680	0.1125
...
8875	-34.7580	185.3760	-0.1875
8876	-153.2024	58.9240	-2.6000
8877	225.6000	480.0000	0.4700
8878	9.5340	34.0500	0.2800
8879	92.5056	192.7200	0.4800

8880 rows × 3 columns

In [289...

```
superstore_order.groupby(["Category", "Sub-Category"]).agg(mean_profit_ratio = ("Profit Rat
```

Out[289...

		mean_profit_ratio
Category	Sub-Category	
Furniture	Bookcases	-0.127756
	Chairs	0.045028
	Furnishings	0.140782
	Tables	-0.147916
Office Supplies	Appliances	-0.145513
	Art	0.251678
	Binders	-0.191641
	Envelopes	0.421913
	Fasteners	0.301157
	Labels	0.429984
	Paper	0.425586
	Storage	0.092382
	Supplies	0.104970
Technology	Accessories	0.219012
	Copiers	0.317826
	Machines	-0.059535
	Phones	0.118926

[Q8] Describe an information that the result above conveys?

Ans: **The results above show the average profit ratio for each subcategory in each category, with positive values indicating profit and negative values indicating loss.**

Pivot and Melt

Pivot

In [290...

```
superstore_order.pivot_table(index="State", columns="Ship Mode", values="Order ID", aggfunc
```

Out[290...

Ship Mode	First Class	Same Day	Second Class	Standard Class
State				
Alabama	9.0	1.0	18.0	30.0
Arizona	42.0	15.0	22.0	123.0
Arkansas	10.0	2.0	8.0	35.0
California	302.0	106.0	346.0	1000.0
Colorado	43.0	5.0	32.0	95.0
Connecticut	19.0	8.0	11.0	39.0
Delaware	16.0	2.0	13.0	55.0
District of Columbia	0.0	0.0	3.0	7.0
Florida	47.0	25.0	57.0	210.0
Georgia	19.0	15.0	31.0	108.0

In [291...

```
pivot_table_result = superstore_order.pivot_table(index="State", columns="Ship Mode", value
pivot_table_result
```

Ship Mode	First Class	Same Day	Second Class	Standard Class
State				
Alabama	9.0	1.0	18.0	30.0
Arizona	42.0	15.0	22.0	123.0
Arkansas	10.0	2.0	8.0	35.0
California	302.0	106.0	346.0	1000.0
Colorado	43.0	5.0	32.0	95.0
Connecticut	19.0	8.0	11.0	39.0
Delaware	16.0	2.0	13.0	55.0
District of Columbia	0.0	0.0	3.0	7.0
Florida	47.0	25.0	57.0	210.0
Georgia	19.0	15.0	31.0	108.0
Idaho	3.0	0.0	2.0	13.0
Illinois	58.0	24.0	96.0	249.0
Indiana	13.0	3.0	30.0	79.0
Iowa	1.0	1.0	4.0	17.0
Kansas	6.0	1.0	2.0	15.0
Kentucky	12.0	5.0	49.0	62.0
Louisiana	7.0	2.0	14.0	15.0
Maine	0.0	0.0	0.0	5.0
Maryland	18.0	7.0	12.0	63.0
Massachusetts	14.0	4.0	35.0	71.0
Michigan	20.0	16.0	43.0	151.0
Minnesota	9.0	4.0	13.0	59.0
Mississippi	3.0	4.0	7.0	36.0
Missouri	7.0	2.0	20.0	24.0
Montana	1.0	1.0	0.0	13.0
Nebraska	6.0	3.0	6.0	20.0
Nevada	4.0	1.0	12.0	17.0
New Hampshire	2.0	0.0	10.0	13.0
New Jersey	5.0	1.0	20.0	87.0
New Mexico	1.0	0.0	9.0	22.0
New York	155.0	57.0	183.0	606.0
North Carolina	36.0	14.0	40.0	139.0
North Dakota	0.0	0.0	5.0	2.0
Ohio	66.0	47.0	84.0	199.0

	Ship Mode	First Class	Same Day	Second Class	Standard Class
State					
	Oklahoma	5.0	6.0	7.0	44.0
	Oregon	20.0	0.0	15.0	81.0
	Pennsylvania	103.0	9.0	78.0	341.0
	Rhode Island	16.0	0.0	21.0	16.0
	South Carolina	3.0	5.0	18.0	16.0
	South Dakota	2.0	0.0	0.0	9.0
	Tennessee	21.0	2.0	24.0	118.0
	Texas	125.0	37.0	161.0	537.0
	Utah	4.0	2.0	19.0	28.0
	Vermont	0.0	0.0	1.0	2.0
	Virginia	39.0	4.0	33.0	115.0
	Washington	56.0	34.0	97.0	265.0
	West Virginia	0.0	0.0	0.0	3.0
	Wisconsin	12.0	3.0	10.0	66.0
	Wyoming	0.0	0.0	0.0	1.0

Melt

In [292...

```
melted_result = pd.melt(pivot_table_result.reset_index(), id_vars=["State"], var_name="Ship Mode", value_name="Order Count")
melted_result
```

Out[292...

	State	Ship Mode	Order Count
0	Alabama	First Class	9.0
1	Arizona	First Class	42.0
2	Arkansas	First Class	10.0
3	California	First Class	302.0
4	Colorado	First Class	43.0
...
191	Virginia	Standard Class	115.0
192	Washington	Standard Class	265.0
193	West Virginia	Standard Class	3.0
194	Wisconsin	Standard Class	66.0
195	Wyoming	Standard Class	1.0

196 rows × 3 columns

[Q9] What is the advantage of using `melt` ?

Ans: Using `melt()` helps simplify data manipulation by turning wide-format data into a more flexible, readable, and analyzable long format. It's especially useful for tasks like grouping, summarizing, and visualizing data efficiently.

[Q10] From the `superstore_order`, display the ascending order considering values in the 'Profit' column to group the 'Category'.

I have provided 2 answers for [Q10] , since this question is confusing.

In [302... `superstore_order.pivot_table(index="Category", values="Profit", aggfunc="sum").sort_values(`

Out[302...

Profit	
Category	
Furniture	16858.5619
Office Supplies	105827.0238
Technology	133410.4932

In [303... `superstore_order.sort_values(by='Profit', ascending=True).groupby('Category').head(5)`
Display the 5 rows with the lowest profit in each group.

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	
7772	7773	CA-2016-108196	25/11/2016	02/12/2016	Standard Class	CS-12505	Cindy Stewart	Consumer	United States	
683	684	US-2017-168116	04/11/2017	04/11/2017	Same Day	GT-14635	Grant Thornton	Corporate	United States	
3011	3012	CA-2017-134845	17/04/2017	23/04/2017	Standard Class	SR-20425	Sharelle Roach	Home Office	United States	
4991	4992	US-2017-122714	07/12/2017	13/12/2017	Standard Class	HG-14965	Henry Goldwyn	Corporate	United States	
3151	3152	CA-2015-147830	15/12/2015	18/12/2015	First Class	NF-18385	Natalie Fritzier	Consumer	United States	
5310	5311	CA-2017-131254	19/11/2017	21/11/2017	First Class	NC-18415	Nathan Cano	Consumer	United States	
1199	1200	CA-2016-130946	08/04/2016	12/04/2016	Standard Class	ZC-21910	Zuschuss Carroll	Consumer	United States	
2697	2698	CA-2014-145317	18/03/2014	23/03/2014	Standard Class	SM-20320	Sean Miller	Home Office	United States	J
27	28	US-2015-150630	17/09/2015	21/09/2015	Standard Class	TB-21520	Tracy Blumstein	Consumer	United States	P
3324	3325	CA-2014-165309	11/11/2014	15/11/2014	Standard Class	KD-16270	Karen Daniels	Consumer	United States	
2928	2929	US-2017-120390	19/10/2017	26/10/2017	Standard Class	TH-21550	Tracy Hopkins	Home Office	United States	
5320	5321	US-2017-162558	02/10/2017	05/10/2017	First Class	Dp-13240	Dean percer	Home Office	United States	
463	464	CA-2016-109869	22/04/2016	29/04/2016	Standard Class	TN-21040	Tanja Norvell	Home Office	United States	
1369	1370	US-2015-103471	24/12/2015	28/12/2015	Standard Class	JR-15670	Jim Radford	Consumer	United States	

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country
6639	6640	CA-2014-108609	31/10/2014	02/11/2014	Second Class	AJ-10780	Anthony Jacobs	Corporate	United States

15 rows × 23 columns

[Q11] Create a new column that calculates the total price (sale*quantity) before discount then group by 'product id' and 'category', then show the mean of the total price

In [306...

```
#enter your code here
superstore_order["Total Price Before Discount"] = superstore_order["Sales"]*superstore_order["Quantity"]
superstore_order.groupby(["Product ID", "Category"]).agg(mean_total_price_bf_discounting =
```

Out[306...

mean_total_price_bf_discounting		
Product ID	Category	
FUR-BO-10000112	Furniture	7426.566000
FUR-BO-10000330	Furniture	1258.192000
FUR-BO-10000362	Furniture	1726.898000
FUR-BO-10000468	Furniture	426.532400
FUR-BO-10000711	Furniture	3194.100000
...
TEC-PH-10004912	Technology	747.320000
TEC-PH-10004922	Technology	673.249500
TEC-PH-10004924	Technology	57.149333
TEC-PH-10004959	Technology	412.009000
TEC-PH-10004977	Technology	2441.475429

1846 rows × 1 columns

[Q12] Complete the function to apply `ratio` column that calculates from `First Class` and `Standard Class` columns on `pivot_table_result`

In [295...

```
# [Q12] Complete the function to apply `ratio` column that calculates from `First Class` and `Standard Class` columns on `pivot_table_result`

# function to transform the ratio
def get_class_ratio(row):

    # get the first class column
    first_class = row['First Class']

    # get the standard class column
    standard_class = row['Standard Class']

    # calculate the ratio
    ratio = first_class / standard_class
```

```

    return ratio

pivot_table_result["ratio"] = pivot_table_result.apply(get_class_ratio, axis=1)

pivot_table_result.head()

```

Out[295...

Ship Mode	First Class	Same Day	Second Class	Standard Class	ratio
State					
Alabama	9.0	1.0	18.0	30.0	0.300000
Arizona	42.0	15.0	22.0	123.0	0.341463
Arkansas	10.0	2.0	8.0	35.0	0.285714
California	302.0	106.0	346.0	1000.0	0.302000
Colorado	43.0	5.0	32.0	95.0	0.452632

[Q13] After complete Q12, What does the `apply` function do?

Ans: **`apply` function in Pandas is used to apply function along the rows or columns of a DataFrame**

[Q14] Create a new column(`short_ratio`) that works the same as Q12 but with `lambda` function

```

In [ ]: # [Q14] Create a new column(`short_ratio`) that works the same as Q12 but with `lambda` fun

pivot_table_result["short_ratio"] = pivot_table_result.apply(lambda row: row['First Class']

pivot_table_result.head()

```

Out[]:

Ship Mode	First Class	Same Day	Second Class	Standard Class	ratio	short_ratio
State						
Alabama	9.0	1.0	18.0	30.0	0.300000	0.300000
Arizona	42.0	15.0	22.0	123.0	0.341463	0.341463
Arkansas	10.0	2.0	8.0	35.0	0.285714	0.285714
California	302.0	106.0	346.0	1000.0	0.302000	0.302000
Colorado	43.0	5.0	32.0	95.0	0.452632	0.452632

[Q15] What is the difference between using `function` in `apply` and `lambda` function? give 2 examples use case.

Ans: A **`function`** is generally used when there is a need to reuse logic, handle complex calculations, or perform multiple operations, but it requires a prior definition. In contrast, a **`lambda` function** is perfect for simple and concise logic, especially when the operation is a one-time task that doesn't need to be reused or require a prior definition.

Example use case :

In [317...

```

df1 = pd.read_csv('./sources/lab/Superstore/superstore_order.csv')

```

```
In [318... df1['Ship Mode'].unique()
```

```
Out[318... array(['Second Class', 'Standard Class', 'First Class', 'Same Day'],  
      dtype=object)
```

```
In [319... # using apply regular function to define the shipping rate by shipping mode (ignore a dista  
def calculate_shipping_rate(row):  
    if row['Ship Mode'] == 'Same Day':  
        shipping_rate = 100.00  
    elif row['Ship Mode'] == 'First Class':  
        shipping_rate = 80.00  
    elif row['Ship Mode'] == 'Second Class':  
        shipping_rate = 50.00  
    else:  
        shipping_rate = 30.00  
  
    return shipping_rate  
  
df1['Ship Rate'] = df1.apply(calculate_shipping_rate, axis=1)  
  
df1[['Ship Mode', 'Ship Rate']]
```

```
Out[319...      Ship Mode  Ship Rate
```

0	Second Class	50.0
1	Second Class	50.0
2	Second Class	50.0
3	Standard Class	30.0
4	Standard Class	30.0
...
8875	Standard Class	30.0
8876	Standard Class	30.0
8877	Standard Class	30.0
8878	Standard Class	30.0
8879	First Class	80.0

8880 rows × 2 columns

```
In [320... # using apply lambda to calculate total price  
superstore_order['Total Price'] = superstore_order.apply(lambda row : row['Total Price Befc  
  
superstore_order[['Total Price Before Discount', 'Discount', 'Total Price']]
```

Out[320...

	Total Price Before Discount	Discount	Total Price
0	523.9200	0.00	523.920000
1	2195.8200	0.00	2195.820000
2	29.2400	0.00	29.240000
3	4787.8875	0.45	2633.338125
4	44.7360	0.20	35.788800
...
8875	370.7520	0.20	296.601600
8876	58.9240	0.80	11.784800
8877	1920.0000	0.00	1920.000000
8878	102.1500	0.00	102.150000
8879	2119.9200	0.00	2119.920000

8880 rows × 3 columns